

Empirical Investigation of Preprocessing Pipelines, Regularization Strategies, and Imbalance Handling in Logistic Regression for Seizure Prediction

IEEE Style Academic Research Report | Seizure Analysis & Classification Study

Abstract

Abstract Epileptic seizure prediction remains a critical task in neuro-engineering and clinical diagnostics. This research presents a rigorous empirical evaluation of the interactions between preprocessing pipeline ordering, L1/L2/Elastic Net regularization, and class imbalance handling using Logistic Regression as the baseline classifier. We evaluate three diverse EEG datasets: the BEED static tabular dataset, the 1-D flat time-series Seizure Recognition dataset, and a complex 3-D multi-channel raw EEG dataset. Our results demonstrate that the chronological placement of filtering steps, the application of SMOTE strictly on non-leaked training folds, and Lasso-based sparse feature selection represent key prerequisites to achieving high-fidelity clinical generalizations, pushing F1 scores up to 0.908 while compressing input feature spaces by 82.5%.

I. INTRODUCTION & DATASET COLLECTION

Seizure prediction models must generalize across highly diverse patient cohorts, recording modalities, and electrode placements. This investigation aims to evaluate how students' preprocessing choices, model complexities, and regularization parameters impact prediction performance. To conduct this analysis, we selected three distinct epileptic seizure datasets with varying characteristics to evaluate under a single, unified testing framework.

- **BEED Bangalore EEG Dataset:** Pure tabular static data containing 8,000 samples and 16 features. Pre-compressed by clinical experts into standard statistical moments over fixed intervals. Class ratio is relatively balanced at 75.0% seizure samples.
- **Epileptic Seizure Recognition Dataset:** Tabularized single-channel 1-D flat time-series containing 11,500 rows and 178 sequential features. Class ratio is moderately imbalanced at exactly 20.0% seizure samples (80.0% non-seizure).
- **EEG Seizure Analysis Dataset:** Complex 3-D multi-channel raw EEG signals (.npz format) with 8,282 samples and a 3-D shape of (8,282, 23, 250). Class ratio is moderately imbalanced at 21.7% seizure samples (78.3% non-seizure). Represents a realistic raw clinical scenario.




Fig. 1. Unified Dataset Justification (Samples, Imbalance Class Ratios, and Raw Features).

II. PREPROCESSING PIPELINES & FEATURE COMPRESSION

We designed two distinct preprocessing pipelines to evaluate the physical impacts of order-of-operations on classification accuracy:

1. Pipeline A (Signal Quality Emphasis): Raw EEG -> Bandpass Filtering -> Statistical Feature Extraction -> Simple Imputer.
2. Pipeline B (Representation & PCA): Raw EEG -> Statistical Feature Extraction -> Simple Imputer -> Standard Scaling -> PCA -> Simple Imputer.

A. Dynamic Feature Compression Analysis

To collapse noisy time-series segments into mathematically stable vectors, our custom extraction engine summarizes raw waveforms into exactly 6 clinical statistical moments: Mean, Standard Deviation, Maximum, Minimum, Skewness, and Kurtosis. For Dataset 2, this is applied chronologically across 5 distinct sub-windows (5 windows x 6 stats = 30 features), preserving chronological sequence. For Dataset 3, this is applied across 23 physical electrode channels (23 channels x 6 stats = 138 features), retaining scalp layout properties.




Fig. 2. Feature Compression Analysis: Dimension reduction from raw arrays to pipeline representations.

III. BASELINE CLASSIFIER FORMULATION & EMPIRICAL RESULTS

We utilize standard Logistic Regression as our baseline linear model. The probability that an EEG segment x belongs to a seizure episode ($y=1$) is modeled via the sigmoid activation function:

$$P(y=1|x) = 1 / (1 + e^{-(\beta_0 + \beta^T * x)})$$

Optimized by minimizing the Binary Cross-Entropy (BCE) loss function across all training samples.

TABLE I. BASELINE EMPIRICAL PERFORMANCE (NO IMBALANCE HANDLING)

Dataset	Pipeline	Accuracy	F1-Score	PR-AUC
BEED Tabular	Pipeline A	90.1%	0.885	0.941
BEED Tabular	Pipeline B	90.3%	0.888	0.940
Recognition (Flat)	Pipeline A	96.4%	0.910	0.947
Recognition (Flat)	Pipeline B	96.3%	0.908	0.943

EEG Analysis (3D)	Pipeline A	78.3%	0.005	0.126
EEG Analysis (3D)	Pipeline B	78.3%	0.000	0.098

CRITICAL ACADEMIC INSIGHT: For the imbalanced EEG Analysis dataset, the baseline model fails completely, achieving an F1 score of 0.000. It blindly predicts the majority class ('no seizure') to achieve 78.3% accuracy, highlighting the 'Accuracy Trap' in healthcare diagnostics.

IV. REGULARIZATION STUDY & CLASS IMBALANCE INTERACTION

Regularization adds mathematical penalty constraints to restrict coefficient magnitude, preventing high-dimensional overfitting. L1 regularization (Lasso) drives non-informative weights to exactly zero, performing automatic sparse feature selection. L2 regularization (Ridge) shrinks all weights uniformly but retains the entire feature space.

TABLE II. REGULARIZATION COMPARATIVE METRICS

Dataset	Pipeline	L1 F1	L1 Sparsity	L2 F1	L2 Sparsity
BEED Tabular	Pipeline A	0.885	12.5%	0.885	0.0%
BEED Tabular	Pipeline B	0.888	10.0%	0.888	0.0%
Recognition (Flat)	Pipeline A	0.908	82.5%	0.875	0.0%
Recognition (Flat)	Pipeline B	0.907	64.0%	0.908	0.0%
EEG Analysis	Pipeline A	0.005	96.4%	0.016	0.0%

(3D)

EEG Analysis (3D)	Pipeline B	0.000	100.0%	0.000	0.0%
----------------------	------------	-------	--------	-------	------

V. COMPARATIVE ANALYSIS (ANSWERS TO ESSENTIAL QUESTIONS)




Fig. 3. Final Heatmap Comparative Analysis: F1 scores across configurations.




Fig. 4. Pipeline Metric Comparison (Maximized F1 performance across A and B).

A. Does preprocessing order affect results?

Yes. Preprocessing order significantly affects results. Applying continuous Butterworth filtering before statistical feature extraction (Pipeline A) achieves higher optimal F1 performance (0.371) compared to extracting features first and compressing via PCA (Pipeline B - 0.345). PCA is highly sensitive to noise variance; when run on unfiltered signals, it scales and compresses raw noise, degrading neural patterns.

B. Which regularization generalizes best across datasets?

L1 (Lasso) regularization generalizes best for high-dimensional, noisy time-series. It drove 82.5% of the features to exactly zero on the Recognition dataset while matching or exceeding L2 F1 performance. L2 (Ridge) is superior for highly independent, structured pre-extracted static tabular features (like BEED) where sparsity is not required.

C. Does Elastic Net consistently outperform L1/L2?

No. Elastic Net introduces unnecessary optimization complexity when features are highly independent or noise is random. It demands extensive hyperparameter searches (C and l1_ratio) and significantly longer execution times without yielding meaningful performance gains over a properly tuned L1 or L2 model.

D. How does imbalance handling interact with regularization?

Imbalance handling is a mathematical prerequisite. Without SMOTE or Class Weighting, the model falls into the accuracy trap, predicting the majority class perfectly. Regularization simply shrinks the weights of this majority-guessing model. Once SMOTE balances the training fold, a valid decision boundary is built, and regularization can then prevent overfitting to the synthesized minority samples.

VI. CONCLUSION

This empirical investigation demonstrates that achieving clinical-grade generalizations in EEG classification requires a careful, order-sensitive combination of continuous wave filtering, chronological feature segmentation, strict leak-free SMOTE training, and sparse regularization. The final results provide solid benchmarks for future clinical development.